

Project report

Course teacher

Professor Vincenzo Caglioti

Group

Project supervisors

Professor Giacomo Boracchi

Professor Diego Carrera

Alex Perugini

876359

Vincenzo Scotti

875505

Marco Re

873564

Problem formulation

- The proposed project consisted in the implementation of a convolutional neural network for the classification of sea lions using pictures extracted from aerial images
- Original Kaggle competition aim was to count the population of sea lions, dividing them into 5 classes: adult male, adult female, subadult male, juvenile and puppy
- The requirement assigned to our group for the project was to provide a classifier able to distinguish between sea lion and background
- Dataset provided by Kaggle, composed of images and relative ground truth



Sample image from the dataset

State of the art

- Image classification is a fundamental problem and at the basis of other computer vision tasks
- Traditionally handcrafted features were first extracted from images using feature descriptors and then used to train the classifier
- In recent years, deep learning exploiting multiple layers of nonlinear information processing have become the main trend
- CNNs are the leading architecture for most image recognition, classification and detection tasks
- Fully CNN exceeded state of the art for semantic segmentation and object detection tasks

Solution approach

- Very large images which can't be used as a whole to train the network
 - Extracted patches from the dataset and used them to train the classifier
- Each image contains more background patches than sea lions ones, this implies that the dataset is very unbalanced
 - Applied augmentation to sea lions patches to partially overcome this problem, increasing a lot the dimension of the original dataset
- Original problem was to count sea lions in the provided image
 - Extended our work to provide the heatmap of the image and to estimate the number of sea lions in it

Implementation: Dataset creation

- Dataset composed of 947 images
 - First 750 used for training
 - Remaining used for testing
- Extraction performed through the absolute difference between the original image and the ground truth
 - Gathering of coordinates and classes of all sea lions in the image
- Background patches extracted by a sliding window over the image cutting all the patches not intersecting a sea lion one
- 96x96 patches



Detail of a sample image and its related ground truth

Implementation: Dataset creation

- Image difference and subsequent blob detection for sea lions coordinates
 - train: 40411
 - validation: 9668
 - test: 13539
- Sliding window and deletion of patches containing sea lions from previous section for background coordinates
 - train: 1129863
 - validation: 9668
 - test: 277390

Implementation: Data augmentation

- Dataset highly unbalanced
 - ◆ too much background patches
 - ◆ not enough sea lions
- Classifier becomes rotation, scale and position invariant so it allows to achieve more robust classification
 - Rotation of a random degree in the range 0 to 360 degrees
 - Flipping both vertically and horizontally
 - Shifting of a maximum of 10 pixels, both vertically and horizontally
 - Zooming of a maximum of 20 pixels

Implementation: Data augmentation samples



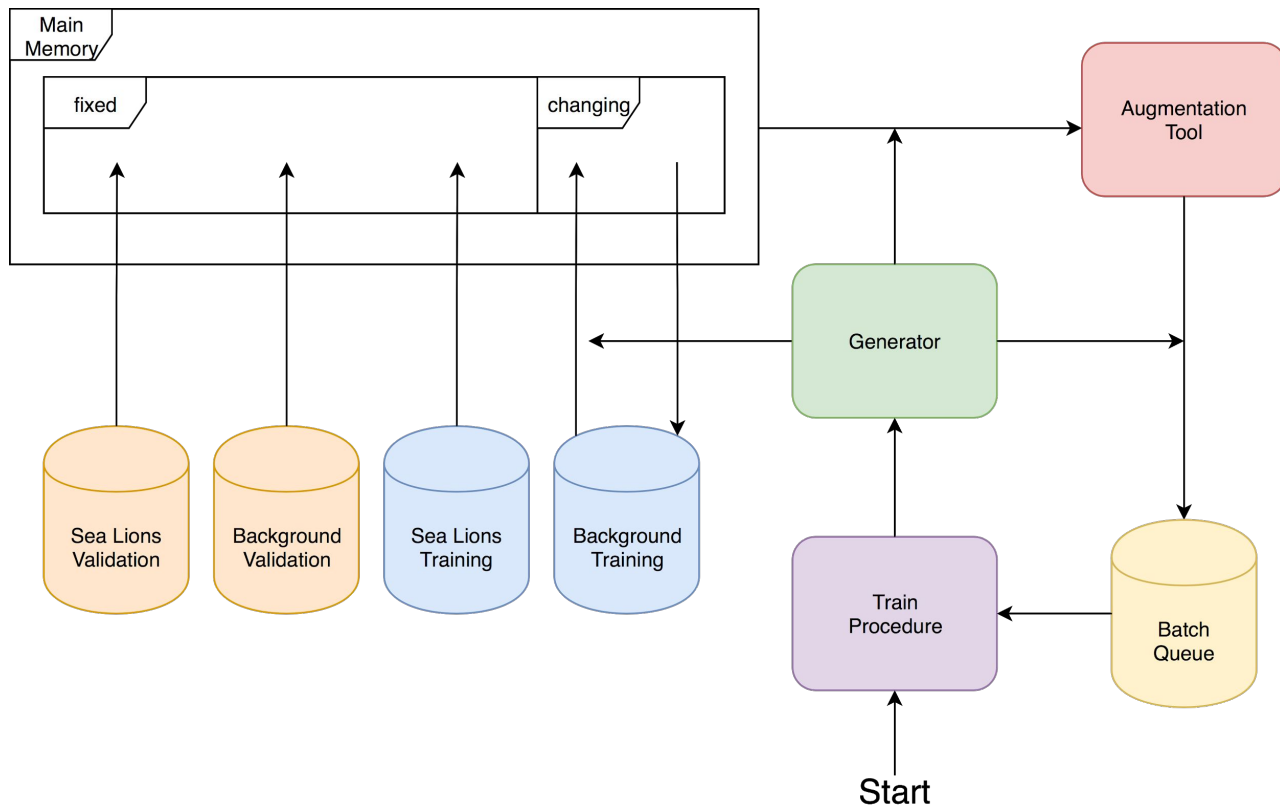
- New sea lions patches generated from original one
- Dataset more balanced




Implementation: Batch generator

- Too little main memory and device memory on the GPU and changing train set led to necessity of efficient memory management
- With batch generator data are moved to the training model only when needed
- Batch generator leverages Keras queue system to avoid waiting for transformations
 - Each batch is generated as the previous are being processed
- Fixed size batch avoids busy memory explosion when moving from 8 bit pixel intensity to normalized float 32 representation
- Portions of the data set for both train and validation that fits in memory are still loaded entirely
 - necessary for sea lions to speed up augmentation
 - background patches are periodically flushed and re-sampled instead

Implementation: Batch generator



Implementation: Model



Implementation: Model

- optimizer: Adam, a stochastic optimization method which uses gradients and second moment gradients to perform parameter update
- loss function: binary cross entropy, error function based on class probabilities
- metric: accuracy, evaluation of the model based on how much the classification is accurate

Implementation: Object detection

- Modified the CNN model to get the fully convolutional one
- Final layer of the network modified into a convolutional layer using 2 filters and a 4x4 kernel, softmax activation function
- Same weights and connections of the previous network
- Dropped constraints on the input shape to use the whole image as input
- Output composed of two heatmaps, highlighting respectively pixels belonging to the sea lion and the background classes according to the probabilities predicted by the classifier
- Cubic interpolation to reconstruct image size heatmap

Experimental activity and results:



Development tools

- Jupyter notebooks
 - Allowed to write Python code and run it section by section
- Keras
 - High level neural network API with TensorFlow as backend
- Dataframes
 - Simplified a lot the management of all the informations about the data set

Experimental activity and results:

Balanced dataset

- As simple as possible
 - No augmentation
 - Balanced train set and test set
- Training
 - 100 epochs
 - Learning rate 0.001



Experimental activity and results:


Balanced dataset

	Loss	Accuracy
Train set	0.0768	0.9733
Validation set	0.0891	0.9679

Prediction accuracy	96.60%
AUC	0.9929


Experimental activity and results: Batch generator and augmentation

- Attempt to use entire data set
 - Use of a custom generator to have balanced data set at each epoch of the training by subsampling randomly the background patches at runtime
 - Efficient memory management through custom generator
- Augmentation
 - Handcrafted augmentation tool
- Training
 - 30 epochs
 - Learning rate 0.0005



Experimental activity and results: Batch generator and augmentation

- Poor results
 - Changing dataset at each epoch is time consuming
 - slow augmentation tool
 - hard to identify/distinguish puppies
- Near random classifier performances



Prediction accuracy	94.91%
AUC	0.5864



	Precision	Recall
Sea lions	0.2948	0.0288
Background	0.9516	0.9972

Experimental activity and results:

Keras augmentation and puppies exclusion

- Dataset reloaded from memory every 15 epochs
 - Speeded up the training
- Puppies removed from the dataset
 - For the network it's easier to learn
 - Number of puppies can be estimated from the number of female sea lions detected
- Augmentation
 - Kears built-in one
 - Higher throughput
- Training
 - 60 epochs
 - Learning rate 0.0005


Experimental activity and results: Keras augmentation and puppies exclusion



Experimental activity and results:

Keras augmentation and puppies exclusion


- Very good performances
- ROC almost covers all the area
- AUC higher than the first one



	Precision	Recall
Sea lions	0.6253	0.9768
Background	0.9989	0.9275

Prediction accuracy	97.27%
AUC	0.9964

Experimental activity and results: Keras augmentation and puppies exclusion





	Precision	Recall
Sea lions	0.6253	0.9768
Background	0.9989	0.9275

Experimental activity and results:

Misclassified patches

- Reasonable mistakes
- Some rocks can appear as sea lions due to their shape and size
- Some sea lions are really hard to see and detect



Experimental activity and results:

Object detection and sea lions counting

- Network reshaping into Fully CNN

- Final layer from flat to convolutional
- Entire image evaluation

- Heatmaps

- Detection

- Heatmap thresholding
- Blob detection
- Counting


	True	Predicted
Image 755	212	233
Image 771	149	141
Image 773	391	411
Image 907	93	123

Experimental activity and results: Object detection and sea lions counting

Input




Heatmap




Experimental activity and results: Object detection and sea lions counting

Sea Lions Area



Detected Sea Lions



Experimental activity and results:

Object detection and sea lions counting

- Non-optimal solution
 - Many sea lions are not identified
 - Many background areas (often rocks) are counted as sea lion
 - Blobs are not correctly positioned even in the correctly identified instances
 - Predicted number is an overestimate
 - True and predicted numbers are similar but puppies are not considered in the prediction

Conclusions

- Final network has good performance on the classification of sea lions over background
- Two major enhancements that allowed this high performances
 - Augmentation
 - Management of the data set during the training
- **Augmentation**
 - At first made results worse
 - Too high variety introduced in the data set with respect to the training we could do
 - After memory handling optimization, provided expected improvements
- **Management of the data set during the training**
 - Allowed to speed up the training a lot by optimizing the memory usage
 - Allowed the network to deal with the augmented data set, increasing robustness of the classification

Open problems

- Reintroduction of the puppies in the data set
- Extension from binary classifier to multiclass classifier
- Transfer learning to improve performances
- Object recognition approach without going through the classification task